1. Title

- TradeFlow AI. Track: Reinvent Commerce.
- Small cross-border sellers lose margin negotiating factory quotes alone, or pay an agent to do it.
- One operator. The engine sets the price. A human sends the offer.
- Prototype. No factory has closed a deal through this.
- Public repo: https://github.com/Fu93/tradeflow-ai. License: MIT.

2. Problem

- Cross-border buyers either negotiate alone, often with a language penalty, or pay an agent.
- TradeFlow is a prototype of a third arrangement: one operator, software doing the arithmetic, the person l
- It is not a deployed company and it has not bought from a real factory.

3. Market, revenue, competition

- Target user: a small cross-border seller who sources from factories and negotiates in chat.
- Revenue model, proposed and not earned: a desk subscription, or a share of savings versus the first-rou
- Competitors: LLM negotiation agents that let the model set the price; sourcing agents who take a commis
- Roadmap: a public offline demo, AMD live reads only when a key is set, measured extraction accuracy o
- No market-size number is stated. None was sourced for this pack.

4. Split

- Hybrid desk. The criteria this workflow is written for are Hybrid, Becreatives, and Vibe Generation.
- Solver decides. The decision core does not ask a language model what to pay.
- Model phrases. An optional call can classify a message or draft wording.
- Human sends. The Chrome extension (Manifest V3) shows a card. It cannot press Send.
- Code: src/engine/solver.py, src/engine/multi_solver.py, extension/.

5. Decision core

- A belief about the supplier reserve, then a vectorised Expectimax search over openings and concession c
- Read the recommended price and the breakdown probability off the dashboard you just captured.
- This slide does not state a latency. Use the number on the screen.

6. AMD language layer

- Sample endpoint: https://developer.amd.com.cn/radeon/api/v1.
- Committed log results/real_llm_benchmark_100.json (2026-10-15): model Qwen3.8-Flash-Next, 348 calls
- That file does not store an HTTP status. status: ok from scripts/verify_amd_endpoint.py means that proce
- The default model for a new Stage-2 call is Gemma-4-31B-IT. The committed log above is Qwen3.8-Flash

7. AMD compute layer

- TRADEFLOW_SOLVER_BACKEND selects NumPy (default) or a CuPy/ROCm request.
- results/bench.json is a CPU measurement (is_gpu_active: false).
- The MI300X column is a theoretical roofline lower-bound estimate, not measured.
- Formula: lower_bound_seconds = (working_set_mb * 1024 * 1024) / (5.3e12).
- No AMD GPU measurement is in the repo unless a sibling benchmark JSON says ran_on_amd_gpu: true

8. Evidence, with the baseline named

- Self-play (results/self_play_benchmark.md): engine savings versus that table’s walk-away ceiling run from
- Real-model log, different experiment: Bluffing Hardliner deal rate 61.11%. 0 budget violations in 94 compl
- Held-out intent accuracy on the synthetic set: 63.08% (195 of 500 held out).
- Auction headline is savings versus the first-round best quote. Versus target is shown, including when that
- Headline savings vs first-round best quote: 5.0% (CNY 6,600.00 total)
- Versus target: -5.0% (CNY -6,000.00 total). A negative figure means the settled price is above target.
- Secondary, versus walk-away ceiling: 22.2% (CNY 36,000.00 total)
- Offline comparison results/llm_vs_tradeflow.json: 24 episodes, seed 42, buyer offline_accept_opening. N

9. Limits and roadmap

- cp .env.example .env && docker compose up --build, then open http://localhost:8000. render.yaml is the p
- Placeholder keys are ignored. /health reports offline_demo and the solver backend. AMD live mode stays
- No factory deal. No shipped MI300X timing. No committed Gemma latency. MIT. Public repo.
- Roadmap: keep the human on the send button, log a real Gemma chat only when a key is present, and p
- Working deadline used for this pack: 11:00 London (BST), 18 October 2026. Confirm the form.